

AMD Radeon AI PRO R9700

显卡测试报告

作者：朴赛技术部

目录

一、AMD 驱动和 Rocm 安装	2
二、显卡压力测试	2
三、ImageNet 模型训练	3
四、模型推理性能测试	4
五、模型推理压力测试	5
六、显卡通信带宽测试	11
七、拓扑展示	14
八、Slurm 调用显卡测试	15

配置	型号	数量
处理器	Threadripper PRO 5975WX	2
显卡	AMD Radeon AI PRO R9700	8
内存	三星 RECC DDR4 32GB 3200MHz	16
硬盘	PM893 SATA 960G	1
网卡	云尖 Intel i350-AM4 四口千兆	1
主板	同泰怡 TG659V2 4U	1

amdgpu 安装完毕:

Rocm 已安装完毕:

使用工具: <https://github.com/ROCm/HIP-Examples>

测试完毕，功耗正常，温度正常，无明显异常：

Epoch:	[46431/5005]	Time	0.073 (0.133)	Data	0.000 (0.034)	Loss	6.9099e+00	(6.9147e+00)	Acc@1	0.00	(0.10)	Acc@5	0.00	(0.50)
Epoch:	[46431/5005]	Time	0.072 (0.133)	Data	0.000 (0.034)	Loss	6.9091e+00	(6.9147e+00)	Acc@1	0.00	(0.10)	Acc@5	0.39	(0.50)
Epoch:	[46451/5005]	Time	0.134 (0.133)	Data	0.000 (0.034)	Loss	6.9048e+00	(6.9147e+00)	Acc@1	0.78	(0.10)	Acc@5	1.56	(0.50)
Epoch:	[4661/5005]	Time	0.099 (0.133)	Data	0.000 (0.034)	Loss	6.9044e+00	(6.9147e+00)	Acc@1	0.00	(0.10)	Acc@5	0.78	(0.50)
Epoch:	[4671/5005]	Time	0.099 (0.133)	Data	0.000 (0.034)	Loss	6.9087e+00	(6.9147e+00)	Acc@1	0.00	(0.10)	Acc@5	0.39	(0.50)
Epoch:	[4681/5005]	Time	0.093 (0.133)	Data	0.000 (0.034)	Loss	6.9112e+00	(6.9147e+00)	Acc@1	0.39	(0.10)	Acc@5	0.39	(0.50)
Epoch:	[4691/5005]	Time	0.071 (0.133)	Data	0.000 (0.034)	Loss	6.9121e+00	(6.9146e+00)	Acc@1	0.00	(0.10)	Acc@5	0.39	(0.50)
Epoch:	[4701/5005]	Time	0.076 (0.133)	Data	0.000 (0.034)	Loss	6.9106e+00	(6.9146e+00)	Acc@1	0.00	(0.10)	Acc@5	0.00	(0.50)
Epoch:	[4711/5005]	Time	0.121 (0.133)	Data	0.000 (0.034)	Loss	6.9076e+00	(6.9146e+00)	Acc@1	0.00	(0.10)	Acc@5	0.00	(0.50)
Epoch:	[4721/5005]	Time	0.098 (0.133)	Data	0.000 (0.034)	Loss	6.9098e+00	(6.9146e+00)	Acc@1	0.00	(0.10)	Acc@5	0.00	(0.50)
Epoch:	[4731/5005]	Time	0.119 (0.133)	Data	0.000 (0.034)	Loss	6.9094e+00	(6.9146e+00)	Acc@1	0.39	(0.10)	Acc@5	0.78	(0.50)
Epoch:	[4741/5005]	Time	0.107 (0.133)	Data	0.000 (0.034)	Loss	6.9064e+00	(6.9146e+00)	Acc@1	0.00	(0.10)	Acc@5	0.39	(0.50)
Epoch:	[4751/5005]	Time	0.135 (0.133)	Data	0.000 (0.034)	Loss	6.9102e+00	(6.9146e+00)	Acc@1	0.00	(0.10)	Acc@5	0.00	(0.50)
Epoch:	[4761/5005]	Time	0.114 (0.133)	Data	0.000 (0.034)	Loss	6.9068e+00	(6.9146e+00)	Acc@1	0.00	(0.10)	Acc@5	0.78	(0.50)
Epoch:	[4771/5005]	Time	0.127 (0.133)	Data	0.000 (0.034)	Loss	6.9099e+00	(6.9146e+00)	Acc@1	0.00	(0.10)	Acc@5	0.39	(0.50)
Epoch:	[4781/5005]	Time	0.110 (0.133)	Data	0.000 (0.034)	Loss	6.9052e+00	(6.9145e+00)	Acc@1	0.39	(0.10)	Acc@5	1.17	(0.50)
Epoch:	[4791/5005]	Time	0.127 (0.133)	Data	0.000 (0.034)	Loss	6.9102e+00	(6.9145e+00)	Acc@1	0.00	(0.10)	Acc@5	0.78	(0.50)
Epoch:	[4801/5005]	Time	0.090 (0.133)	Data	0.000 (0.034)	Loss	6.9107e+00	(6.9145e+00)	Acc@1	0.00	(0.10)	Acc@5	0.39	(0.50)
Epoch:	[4811/5005]	Time	0.090 (0.133)	Data	0.000 (0.034)	Loss	6.9066e+00	(6.9145e+00)	Acc@1	0.39	(0.10)	Acc@5	1.56	(0.50)
Epoch:	[4821/5005]	Time	0.076 (0.133)	Data	0.000 (0.034)	Loss	6.9086e+00	(6.9145e+00)	Acc@1	0.00	(0.10)	Acc@5	0.39	(0.50)
Epoch:	[4831/5005]	Time	0.072 (0.133)	Data	0.000 (0.034)	Loss	6.9104e+00	(6.9145e+00)	Acc@1	0.39	(0.10)	Acc@5	1.17	(0.50)
Epoch:	[4841/5005]	Time	0.072 (0.133)	Data	0.000 (0.034)	Loss	6.9134e+00	(6.9145e+00)	Acc@1	0.00	(0.10)	Acc@5	0.39	(0.50)
Epoch:	[4851/5005]	Time	0.074 (0.133)	Data	0.000 (0.034)	Loss	6.9053e+00	(6.9145e+00)	Acc@1	0.00	(0.10)	Acc@5	0.78	(0.50)
Epoch:	[4861/5005]	Time	0.075 (0.133)	Data	0.000 (0.034)	Loss	6.9082e+00	(6.9144e+00)	Acc@1	0.00	(0.10)	Acc@5	0.78	(0.50)
Epoch:	[4871/5005]	Time	0.071 (0.133)	Data	0.000 (0.034)	Loss	6.9089e+00	(6.9144e+00)	Acc@1	0.00	(0.10)	Acc@5	0.39	(0.50)
Epoch:	[4881/5005]	Time	0.098 (0.133)	Data	0.000 (0.034)	Loss	6.9077e+00	(6.9144e+00)	Acc@1	0.00	(0.10)	Acc@5	0.00	(0.50)
Epoch:	[4891/5005]	Time	0											

四、模型推理性能测试

使用模型： qwen3.5-122B-A10B-FP8

模型参数量： 122B

对接工具： openweb-ui

测试结果： 单并发场景， 峰值 96tokens/s， 性能良好。

```
(APIServer pid=362331) INFO 05-27 15:30:43 [api_server.py:57] Starting vLLM server on http://0.0.0.0:8000
(APIServer pid=362331) INFO 05-27 15:30:43 [launcher.py:37] Available routes are:
(APIServer pid=362331) INFO 05-27 15:30:43 [launcher.py:40] Route: /openapi.json, Methods: HEAD, GET
(APIServer pid=362331) INFO 05-27 15:30:43 [launcher.py:40] Route: /docs, Methods: HEAD, GET
(APIServer pid=362331) INFO 05-27 15:30:43 [launcher.py:40] Route: /docs/openai-redirect, Methods: HEAD, GET
(APIServer pid=362331) INFO 05-27 15:30:43 [launcher.py:40] Route: /redoc, Methods: HEAD, GET
(APIServer pid=362331) INFO 05-27 15:30:43 [launcher.py:40] Route: /tokenize, Methods: POST
(APIServer pid=362331) INFO 05-27 15:30:43 [launcher.py:40] Route: /detokenize, Methods: POST
(APIServer pid=362331) INFO 05-27 15:30:43 [launcher.py:40] Route: /load, Methods: GET
(APIServer pid=362331) INFO 05-27 15:30:43 [launcher.py:40] Route: /perfmon, Methods: GET
(APIServer pid=362331) INFO 05-27 15:30:43 [launcher.py:40] Route: /health, Methods: GET
(APIServer pid=362331) INFO 05-27 15:30:43 [launcher.py:40] Route: /metrics, Methods: GET
(APIServer pid=362331) INFO 05-27 15:30:43 [launcher.py:40] Route: /v1/models, Methods: GET
(APIServer pid=362331) INFO 05-27 15:30:43 [launcher.py:40] Route: /ping, Methods: GET
(APIServer pid=362331) INFO 05-27 15:30:43 [launcher.py:40] Route: /ping, Methods: POST
(APIServer pid=362331) INFO 05-27 15:30:43 [launcher.py:40] Route: /v1/chat/completions, Methods: POST
(APIServer pid=362331) INFO 05-27 15:30:43 [launcher.py:40] Route: /v1/chat/completions/batch, Methods: POST
(APIServer pid=362331) INFO 05-27 15:30:43 [launcher.py:40] Route: /v1/responses, Methods: POST
(APIServer pid=362331) INFO 05-27 15:30:43 [launcher.py:40] Route: /v1/responses/{response_id}/cancel, Methods: GET
(APIServer pid=362331) INFO 05-27 15:30:43 [launcher.py:40] Route: /v1/completions, Methods: POST
(APIServer pid=362331) INFO 05-27 15:30:43 [launcher.py:40] Route: /v1/messages, Methods: POST
(APIServer pid=362331) INFO 05-27 15:30:43 [launcher.py:40] Route: /v1/messages/count_tokens, Methods: POST
(APIServer pid=362331) INFO 05-27 15:30:43 [launcher.py:40] Route: /inference/v1/generate, Methods: POST
(APIServer pid=362331) INFO 05-27 15:30:43 [launcher.py:40] Route: /scale_elastic_ep, Methods: POST
(APIServer pid=362331) INFO 05-27 15:30:43 [launcher.py:40] Route: /is_scaling_elastic_ep, Methods: POST
(APIServer pid=362331) INFO 05-27 15:30:43 [launcher.py:40] Route: /generate_scoring, Methods: POST
(APIServer pid=362331) INFO 05-27 15:30:43 [launcher.py:40] Route: /v1/chat/completions/render, Methods: POST
(APIServer pid=362331) INFO 05-27 15:30:43 [launcher.py:40] Route: /v1/completions/render, Methods: POST
(APIServer pid=362331) INFO 05-27 15:30:43 [launcher.py:40] Started server process [362331]
(APIServer pid=362331) INFO 05-27 15:30:43 [launcher.py:40] Waiting for application startup.
(APIServer pid=362331) INFO 05-27 15:30:43 [launcher.py:40] Application startup complete.
(APIServer pid=362331) INFO 05-27 15:30:43 [launcher.py:40] 192.168.12.1:59069 - "GET /v1/models HTTP/1.1" 200 OK
(APIServer pid=362331) INFO 05-27 15:30:43 [launcher.py:40] 192.168.12.1:60005 - "GET /v1/models HTTP/1.1" 200 OK
(APIServer pid=362331) INFO 05-27 15:30:43 [launcher.py:40] 192.168.12.1:52007 - "GET /v1/models HTTP/1.1" 200 OK
(APIServer pid=362331) INFO 05-27 15:30:43 [launcher.py:40] 192.168.12.1:52042 - "GET /v1/models HTTP/1.1" 200 OK
(APIServer pid=362331) INFO 05-27 15:30:43 [launcher.py:40] 192.168.12.1:52105 - "POST /v1/chat/completions HTTP/1.1" 200 OK
(Worker_T0 pid=362922) WARNING 05-27 16:03:16 [jit_monitor.py:103] Triton kernel JIT compilation during inference: _zero_kv_blocks_kernel. This causes a latency spike; consider extending warmup to cover this shape/config.
(Worker_T0 pid=362922) WARNING 05-27 16:03:16 [jit_monitor.py:103] Triton kernel JIT compilation during inference: _compute_slot_mapping_kernel. This causes a latency spike; consider extending warmup to cover this shape/config.
(Worker_T0 pid=362922) WARNING 05-27 16:03:17 [jit_monitor.py:103] Triton kernel JIT compilation during inference: _causal_conv2d_fwd_kernel. This causes a latency spike; consider extending warmup to cover this shape/config.
(Worker_T0 pid=362922) WARNING 05-27 16:03:17 [jit_monitor.py:103] Triton kernel JIT compilation during inference: _fused_post_conv_kernel. This causes a latency spike; consider extending warmup to cover this shape/config.
(Worker_T0 pid=362922) WARNING 05-27 16:03:18 [jit_monitor.py:103] Triton kernel JIT compilation during inference: _fwd_kernel. This causes a latency spike; consider extending warmup to cover this shape/config.
(APIServer pid=362331) INFO 05-27 16:03:21 [loggers.py:271] Engine 000: Avg prompt throughput: 1.1 tokens/s, Avg generation throughput: 0.1 tokens/s, Running: 1 reqs, Waiting: 0 reqs, GPU KV cache usage: 0.2%, Prefix cache hit rate: 0.0%
(APIServer pid=362331) INFO 05-27 16:03:31 [loggers.py:271] Engine 000: Avg prompt throughput: 35.6 tokens/s, Avg generation throughput: 58.2 tokens/s, Running: 1 reqs, Waiting: 0 reqs, GPU KV cache usage: 0.3%, Prefix cache hit rate: 0.0%
(APIServer pid=362331) INFO 05-27 16:03:41 [loggers.py:271] Engine 000: Avg prompt throughput: 4.2 tokens/s, Avg generation throughput: 76.8 tokens/s, Running: 2 reqs, Waiting: 0 reqs, GPU KV cache usage: 0.5%, Prefix cache hit rate: 0.0%
(APIServer pid=362331) INFO 05-27 16:03:51 [loggers.py:271] Engine 000: Avg prompt throughput: 0.0 tokens/s, Avg generation throughput: 96.0 tokens/s, Running: 2 reqs, Waiting: 0 reqs, GPU KV cache usage: 0.6%, Prefix cache hit rate: 0.0%
(APIServer pid=362331) INFO 05-27 16:04:01 [loggers.py:271] Engine 000: Avg prompt throughput: 42.9 tokens/s, Avg generation throughput: 94.8 tokens/s, Running: 2 reqs, Waiting: 0 reqs, GPU KV cache usage: 0.6%, Prefix cache hit rate: 0.0%
(APIServer pid=362331) INFO 05-27 16:04:11 [loggers.py:271] Engine 000: Avg prompt throughput: 0.0 tokens/s, Avg generation throughput: 95.4 tokens/s, Running: 2 reqs, Waiting: 0 reqs, GPU KV cache usage: 0.7%, Prefix cache hit rate: 0.0%
(APIServer pid=362331) INFO 05-27 16:04:21 [loggers.py:271] Engine 000: Avg prompt throughput: 0.0 tokens/s, Avg generation throughput: 95.4 tokens/s, Running: 2 reqs, Waiting: 0 reqs, GPU KV cache usage: 0.7%, Prefix cache hit rate: 0.0%
(APIServer pid=362331) INFO 05-27 16:04:31 [loggers.py:271] Engine 000: Avg prompt throughput: 0.0 tokens/s, Avg generation throughput: 95.4 tokens/s, Running: 2 reqs, Waiting: 0 reqs, GPU KV cache usage: 0.7%, Prefix cache hit rate: 0.0%
```


五、模型推理压力测试

单并发压力测试: 55.56 tokens/s 平均带宽: 55.56 tokens/s

10 并发压力测试: 233.95 tokens/s 平均带宽: 23.40 tokens/s

50 并发压力测试: 739.45 tokens/s 平均带宽: 14.79 tokens/s

100 并发压力测试: 1048.41 tokens/s 平均带宽: 10.48 tokens/s

200 并发压力测试: 1193.39 tokens/s 平均带宽: 6 tokens/s

结论：该显卡运行该模型无异常，并发量在 1~30 时，有较为客观的性能表现，适合 50 人左右团队使用。

Performance Test Summary Report

Basic Information:

Model	qwen3.5-122B-A10B-FP8
Test Dataset	random
API Type	openai
Total Generated	747,520.0 tokens
Avg Output Rate	787.77 tokens/sec
Total Test Time	948.91 seconds
Output Path	outputs/20260527_162714/qwen3.5-122B-A10B-FP8

Detailed Performance Metrics

					Avg	P99	Avg	P99	Avg	P99	Gen.	Success
Conc.	Rate	Num	RPS	Lat. (s)	Lat. (s)	TTFT(s)	TTFT(s)	TPOT(s)	TPOT(s)	toks/s	Rate	
1	INF	10	0.05	18.422	18.553	0.328	0.402	0.018	0.018	55.56	100.0%	
10	INF	20	0.23	43.645	46.695	1.087	2.663	0.042	0.043	233.95	100.0%	
50	INF	100	0.72	68.215	81.286	4.138	14.486	0.063	0.066	739.45	100.0%	
100	INF	200	1.02	94.794	122.643	9.236	32.704	0.084	0.090	1048.41	100.0%	
200	INF	400	1.17	162.290	235.834	25.984	84.120	0.133	0.157	1193.39	100.0%	

Request Metrics

			Avg In	P99 In	Avg Out	P99 Out	Decoded	Spec.
Conc.	Num		Toks	Toks	Toks	Toks	Tok/Iter	Accept Rate
1	10		1098.3	1167.0	1024.0	1024.0	1.00	0.0%
10	20		1069.2	1180.0	1024.0	1024.0	1.00	0.0%
50	100		1066.2	1198.0	1024.0	1024.0	1.00	0.0%
100	200		1066.1	1188.0	1024.0	1024.0	1.00	0.0%
200	400		1063.6	1196.0	1024.0	1024.0	1.00	0.0%

Best Performance Configuration

Highest RPS Concurrency 200 (1.17 req/sec)

Lowest Latency Concurrency 1 (18.422 seconds)

Performance Recommendations:

- The system seems not to have reached its performance bottleneck, try higher concurrency

```
2026-05-27 16:45:52 - evalscope - INFO: Performance summary saved to: outputs/20260527_162714/qwen3.5-122B-A10B-FP8/performance_summary.txt
```

```
2026-05-27 16:45:53 - evalscope - INFO: HTML report generated: /home/ps/outputs/20260527_162714/qwen3.5-72B-A10B-FP8/perf_report.html
```

单并发

Benchmarking summary:

Key	Value
Time taken for tests (s)	184.321
Number of concurrency	1
Request rate (req/s)	-1
Total requests	10
Succeed requests	10
Failed requests	0
Request throughput (req/s)	0.0543
Average latency (s)	18.4225
Average input tokens per request	1098.3
Output token throughput (tok/s)	55.5552
Total token throughput (tok/s)	115.141
Average time to first token (s)	0.3277
Average time per output token (s)	0.0177
Average inter-token latency (s)	0.0177
Average output tokens per request	1024
Average decoded tokens per iter (tok/iter)	0.999
Approx speculative decoding acceptance rate	0

2026-05-27 16:30:41 - evalscope - INFO:

Percentile results:

Percentiles	TTFT (s)	ITL (s)	TPOT (s)	Latency (s)	Input tokens	Output tokens	Output (tok/s)	Total (tok/s)
10%	0.2854	0.0171	0.0176	18.29	1024	1024	55.2944	111.9738
25%	0.3157	0.0173	0.0176	18.3247	1026	1024	55.3803	112.0937
50%	0.33	0.0176	0.0177	18.4513	1109	1024	55.5477	115.8791
66%	0.3387	0.0178	0.0177	18.4847	1137	1024	55.6307	116.6907
75%	0.339	0.018	0.0177	18.4903	1156	1024	55.8809	117.7177
80%	0.3641	0.018	0.0178	18.519	1160	1024	55.9869	117.8994
90%	0.4024	0.0182	0.0178	18.5529	1167	1024	56.0468	118.5302
95%	0.4024	0.0183	0.0178	18.5529	1167	1024	56.0468	118.5302
98%	0.4024	0.0183	0.0178	18.5529	1167	1024	56.0468	118.5302
99%	0.4024	0.0188	0.0178	18.5529	1167	1024	56.0468	118.5302

10 并发

Benchmarking summary:

Key	Value
Time taken for tests (s)	87.5385
Number of concurrency	10
Request rate (req/s)	-1
Total requests	20
Succeed requests	20
Failed requests	0
Request throughput (req/s)	0.2285
Average latency (s)	43.6451
Average input tokens per request	1069.15
Output token throughput (tok/s)	233.954
Total token throughput (tok/s)	478.224
Average time to first token (s)	1.0875
Average time per output token (s)	0.0416
Average inter-token latency (s)	0.0416
Average output tokens per request	1024
Average decoded tokens per iter (tok/iter)	0.999
Approx speculative decoding acceptance rate	0

2026-05-27 16:32:34 - evalscope - INFO:

Percentile results:

Percentiles	TTFT (s)	ITL (s)	TPOT (s)	Latency (s)	Input tokens	Output tokens	Output (tok/s)	Total (tok/s)
10%	0.5354	0.0329	0.0396	41.2498	1024	1024	22.1599	44.3199
25%	0.6127	0.0411	0.0403	41.8022	1024	1024	22.518	45.6935
50%	0.729	0.0417	0.043	44.6106	1024	1024	24.181	48.362
66%	1.2133	0.0421	0.043	45.2328	1087	1024	24.3925	49.2034
75%	1.6922	0.0423	0.043	45.7166	1142	1024	24.6017	50.0694
80%	1.9318	0.0425	0.043	45.9636	1142	1024	24.7197	50.146
90%	2.418	0.0427	0.0431	46.4531	1177	1024	24.9472	52.5094
95%	2.6628	0.0429	0.0432	46.6953	1180	1024	25.073	53.2054
98%	2.6628	0.0461	0.0432	46.6953	1180	1024	25.073	53.2054
99%	2.6628	0.1191	0.0432	46.6953	1180	1024	25.073	53.2054

50 并发

Benchmarking summary:									
Key	Value								
Time taken for tests (s)	138.481								
Number of concurrency	50								
Request rate (req/s)	-1								
Total requests	100								
Succeed requests	100								
Failed requests	0								
Request throughput (req/s)	0.7221								
Average latency (s)	68.2149								
Average input tokens per request	1066.2								
Output token throughput (tok/s)	739.45								
Total token throughput (tok/s)	1509.37								
Average time to first token (s)	4.1377								
Average time per output token (s)	0.0626								
Average inter-token latency (s)	0.0626								
Average output tokens per request	1024								
Average decoded tokens per iter (tok/iter)	0.999								
Approx speculative decoding acceptance rate	0								
2026-05-27 16:35:23 - evalscope - INFO:									
Percentile results:									
Percentiles	TTFT (s)	ITL (s)	TPOT (s)	Latency (s)	Input tokens	Output tokens	Output (tok/s)	Total (tok/s)	
10%	0.8578	0.054	0.0573	59.9204	1024	1024	13.1506	27.2198	
25%	1.1272	0.0552	0.0608	63.2892	1024	1024	13.9865	28.6902	
50%	1.3	0.0562	0.0643	67.2311	1025	1024	15.253	30.9103	
66%	4.5925	0.0567	0.0648	70.6754	1103	1024	15.7434	32.3467	
75%	7.1342	0.0569	0.065	73.4981	1130	1024	16.2197	33.0305	
80%	8.566	0.0571	0.0651	75.0518	1137	1024	16.5152	33.6583	
90%	11.6304	0.0652	0.0652	78.2847	1152	1024	17.1766	34.828	
95%	13.2737	0.1377	0.0653	80.0153	1180	1024	17.5768	35.812	
98%	14.2504	0.1428	0.0655	81.0059	1193	1024	17.83	37.2606	
99%	14.4859	0.1452	0.0656	81.2857	1198	1024	17.906	37.3335	

100 并发

Benchmarking summary:

Key	Value
Time taken for tests (s)	195.344
Number of concurrency	100
Request rate (req/s)	-1
Total requests	200
Succeed requests	200
Failed requests	0
Request throughput (req/s)	1.0238
Average latency (s)	94.7944
Average input tokens per request	1066.12
Output token throughput (tok/s)	1048.41
Total token throughput (tok/s)	2139.94
Average time to first token (s)	9.2357
Average time per output token (s)	0.0836
Average inter-token latency (s)	0.0836
Average output tokens per request	1024
Average decoded tokens per iter (tok/iter)	0.999
Approx speculative decoding acceptance rate	0

2026-05-27 16:39:17 - evalscope - INFO:

Percentile results:

Percentiles	TTFT (s)	ITL (s)	TPOT (s)	Latency (s)	Input tokens	Output tokens	Output (tok/s)	Total (tok/s)
10%	1.7866	0.0646	0.0727	78.9559	1024	1024	8.9292	18.25
25%	3.0362	0.0665	0.0802	85.453	1024	1024	9.7879	19.951
50%	4.5478	0.0725	0.0873	91.7796	1025	1024	11.1757	22.5771
66%	9.168	0.0749	0.0879	99.2353	1108	1024	11.6156	23.8792
75%	14.6555	0.076	0.088	104.8978	1124	1024	12.0063	24.4709
80%	17.9536	0.0766	0.0881	108.1675	1130	1024	12.2819	24.9767
90%	25.1967	0.1445	0.0882	115.0011	1143	1024	13.0036	26.6825
95%	29.3235	0.1599	0.0883	119.2313	1155	1024	13.575	27.8569
98%	31.7765	0.1653	0.0895	121.6653	1170	1024	13.9281	28.526
99%	32.7037	0.1683	0.0899	122.6426	1188	1024	14.0942	28.9348

200 并发

Benchmarking summary:

Key	Value
Time taken for tests (s)	343.223
Number of concurrency	200
Request rate (req/s)	-1
Total requests	400
Succeed requests	400
Failed requests	0
Request throughput (req/s)	1.1654
Average latency (s)	162.29
Average input tokens per request	1063.63
Output token throughput (tok/s)	1193.39
Total token throughput (tok/s)	2432.97
Average time to first token (s)	25.9842
Average time per output token (s)	0.1332
Average inter-token latency (s)	0.1331
Average output tokens per request	1024
Average decoded tokens per iter (tok/iter)	0.999
Approx speculative decoding acceptance rate	0

2026-05-27 16:45:52 - evalscope - INFO:

Percentile results:

Percentiles	TTFT (s)	ITL (s)	TPOT (s)	Latency (s)	Input tokens	Output tokens	Output (tok/s)	Total (tok/s)
10%	6.0389	0.084	0.102	125.1899	1024	1024	4.8049	9.7705
25%	12.6735	0.087	0.1253	141.0252	1024	1024	5.6717	11.6218
50%	20.2501	0.0995	0.1397	157.1748	1025	1024	6.5203	13.3408
66%	21.2955	0.1766	0.1449	163.4421	1088	1024	6.9521	14.2382
75%	33.5622	0.1931	0.1477	181.1347	1121	1024	7.2632	14.8149
80%	41.5383	0.1961	0.1485	191.5439	1128	1024	7.4757	15.1733
90%	61.2202	0.2045	0.1491	213.6976	1142	1024	8.2082	16.6046
95%	73.4675	0.209	0.1514	225.8444	1158	1024	8.8801	18.1512
98%	81.3998	0.2135	0.1553	233.3338	1186	1024	9.3746	18.8596
99%	84.1195	0.2162	0.1569	235.834	1196	1024	9.5559	19.7099

六、显卡通信带宽测试

注意：测试使用主板插槽为 Pcie4.0

使用命令:./all_gather_perf -g 8 -b 512M -e 1G -i 50m

```
# rccl-tests version 2.17.9-: rccl-headers=22707 rccl-library=22707
# Collective test starting: all_gather_perf
# nThread 1 nGpus 8 minBytes 536870912 maxBytes 1073741824 step: 52428800(bytes) warmup iters: 1 iters: 20 agg iters: 1
validation: 1 graph: 0
#
# Using devices
# Rank 0 Group 0 Pid 471180 on ps device 0 [0000:43:00] AMD Radeon AI PRO R9700
# Rank 1 Group 0 Pid 471180 on ps device 1 [0000:26:00] AMD Radeon AI PRO R9700
# Rank 2 Group 0 Pid 471180 on ps device 2 [0000:03:00] AMD Radeon AI PRO R9700
# Rank 3 Group 0 Pid 471180 on ps device 3 [0000:e3:00] AMD Radeon AI PRO R9700
# Rank 4 Group 0 Pid 471180 on ps device 4 [0000:c3:00] AMD Radeon AI PRO R9700
# Rank 5 Group 0 Pid 471180 on ps device 5 [0000:c6:00] AMD Radeon AI PRO R9700
# Rank 6 Group 0 Pid 471180 on ps device 6 [0000:a3:00] AMD Radeon AI PRO R9700
# Rank 7 Group 0 Pid 471180 on ps device 7 [0000:83:00] AMD Radeon AI PRO R9700
#
#
# out-of-place in-place
# size count type redop root time algbw busbw #wrong time algbw busbw #wrong
# (B) (elements) (us) (GB/s) (GB/s) (us) (GB/s) (GB/s)
536870912 16777216 float none -1 24473.7 21.94 19.19 0 24421.7 21.98 19.24 0
589299712 18415616 float none -1 27175.1 21.69 18.97 0 29249.9 20.15 17.63 0
641728512 20054016 float none -1 29784.9 21.55 18.85 0 29664.9 21.63 18.93 0
694157312 21692416 float none -1 34293.5 20.24 17.71 0 33118.2 20.96 18.34 0
746586112 23330816 float none -1 34247.5 21.80 19.07 0 34235.5 21.81 19.08 0
799014912 24969216 float none -1 36614.1 21.82 19.09 0 36486.6 21.90 19.16 0
851443712 26607616 float none -1 39345.8 21.64 18.94 0 39205.0 21.72 19.00 0
903872512 28246016 float none -1 41834.5 21.61 18.91 0 41687.8 21.68 18.97 0
956301312 29884416 float none -1 44072.8 21.70 18.99 0 43992.1 21.74 19.02 0
1008730112 31522816 float none -1 46145.1 21.86 19.13 0 46160.6 21.85 19.12 0
1061158912 33161216 float none -1 48801.5 21.74 19.03 0 48690.4 21.79 19.07 0
# Out of bounds values : 0 OK
# Avg bus bandwidth : 18.8838
#
# Collective test concluded: all_gather_perf
#
```


使用命令：./all_reduce_perf -g 8 -b 512M -e 1G -i 50m

```
# rccl-tests version 2.17.9-:+ rccl-headers=22707 rccl-library=22707
# Collective test starting: all_reduce_perf
# nThread 1 nGpus 8 minBytes 536870912 maxBytes 1073741824 step: 52428800(bytes) warmup iters: 1 iters: 20 agg iters: 1
validation: 1 graph: 0
#
# Using devices
# Rank 0 Group 0 Pid 471487 on ps device 0 [0000:43:00] AMD Radeon AI PRO R9700
# Rank 1 Group 0 Pid 471487 on ps device 1 [0000:26:00] AMD Radeon AI PRO R9700
# Rank 2 Group 0 Pid 471487 on ps device 2 [0000:03:00] AMD Radeon AI PRO R9700
# Rank 3 Group 0 Pid 471487 on ps device 3 [0000:e3:00] AMD Radeon AI PRO R9700
# Rank 4 Group 0 Pid 471487 on ps device 4 [0000:c3:00] AMD Radeon AI PRO R9700
# Rank 5 Group 0 Pid 471487 on ps device 5 [0000:c6:00] AMD Radeon AI PRO R9700
# Rank 6 Group 0 Pid 471487 on ps device 6 [0000:a3:00] AMD Radeon AI PRO R9700
# Rank 7 Group 0 Pid 471487 on ps device 7 [0000:83:00] AMD Radeon AI PRO R9700
#
#
# out-of-place in-place
# size count type redop root time algbw busbw #wrong time algbw busbw #wrong
# (B) (elements) (us) (GB/s) (GB/s) (us) (GB/s) (GB/s)
#
536870912 134217728 float sum -1 48555.7 11.06 19.35 0 49603.2 10.82 18.94 0
589299712 147324928 float sum -1 55455.8 10.63 18.60 0 54363.2 10.84 18.97 0
641728512 160432128 float sum -1 58157.6 11.03 19.31 0 58174.3 11.03 19.30 0
694157312 173539328 float sum -1 62848.2 11.04 19.33 0 62852.5 11.04 19.33 0
746586112 186646528 float sum -1 67630.8 11.04 19.32 0 67668.6 11.03 19.31 0
799014912 199753728 float sum -1 72397.9 11.04 19.31 0 72305.3 11.05 19.34 0
851443712 212860928 float sum -1 77112.8 11.04 19.32 0 77139.2 11.04 19.32 0
903872512 225968128 float sum -1 81832.2 11.05 19.33 0 81799.4 11.05 19.34 0
956301312 239075328 float sum -1 86558.3 11.05 19.33 0 86654.9 11.04 19.31 0
1008730112 252182528 float sum -1 91362.6 11.04 19.32 0 91297.9 11.05 19.34 0
1061158912 265289728 float sum -1 96115.9 11.04 19.32 0 96123.7 11.04 19.32 0
# Out of bounds values : 0 OK
# Avg bus bandwidth : 19.257
#
# Collective test concluded: all_reduce_perf
#
```

使用命令：./broadcast_perf -g 8 -b 512M -e 1G -i 50m

```
# rccl-tests version 2.17.9-:+ rccl-headers=22707 rccl-library=22707
# Collective test starting: broadcast_perf
# nThread 1 nGpus 8 minBytes 536870912 maxBytes 1073741824 step: 52428800(bytes) warmup iters: 1 iters: 20 agg iters: 1
validation: 1 graph: 0
#
# Using devices
# Rank 0 Group 0 Pid 471797 on ps device 0 [0000:43:00] AMD Radeon AI PRO R9700
# Rank 1 Group 0 Pid 471797 on ps device 1 [0000:26:00] AMD Radeon AI PRO R9700
# Rank 2 Group 0 Pid 471797 on ps device 2 [0000:03:00] AMD Radeon AI PRO R9700
# Rank 3 Group 0 Pid 471797 on ps device 3 [0000:e3:00] AMD Radeon AI PRO R9700
# Rank 4 Group 0 Pid 471797 on ps device 4 [0000:c3:00] AMD Radeon AI PRO R9700
# Rank 5 Group 0 Pid 471797 on ps device 5 [0000:c6:00] AMD Radeon AI PRO R9700
# Rank 6 Group 0 Pid 471797 on ps device 6 [0000:a3:00] AMD Radeon AI PRO R9700
# Rank 7 Group 0 Pid 471797 on ps device 7 [0000:83:00] AMD Radeon AI PRO R9700
#
#
# out-of-place in-place
# size count type redop root time algbw busbw #wrong time algbw busbw #wrong
# (B) (elements) (us) (GB/s) (GB/s) (us) (GB/s) (GB/s)
#
536870912 134217728 float none 0 26794.0 20.04 20.04 0 26812.2 20.02 20.02 0
589299712 147324928 float none 0 29433.4 20.02 20.02 0 30368.8 19.40 19.40 0
641728512 160432128 float none 0 32065.3 20.01 20.01 0 32079.8 20.00 20.00 0
694157312 173539328 float none 0 36304.1 19.12 19.12 0 35520.6 19.54 19.54 0
746586112 186646528 float none 0 37317.3 20.01 20.01 0 37306.9 20.01 20.01 0
799014912 199753728 float none 0 39917.3 20.02 20.02 0 39902.3 20.02 20.02 0
851443712 212860928 float none 0 42519.5 20.02 20.02 0 42533.0 20.02 20.02 0
903872512 225968128 float none 0 45161.0 20.01 20.01 0 45177.7 20.01 20.01 0
956301312 239075328 float none 0 47752.2 20.03 20.03 0 47786.4 20.01 20.01 0
1008730112 252182528 float none 0 50343.3 20.04 20.04 0 50367.3 20.03 20.03 0
1061158912 265289728 float none 0 53014.9 20.02 20.02 0 53000.6 20.02 20.02 0
# Out of bounds values : 0 OK
# Avg bus bandwidth : 19.9287
#
# Collective test concluded: broadcast_perf
#
```


使用命令：./alltoall_perf -g 8 -b 512M -e 1G -i 50m

```
# rccl-tests version 2.17.9-: rccl-headers=22707 rccl-library=22707
# Collective test starting: alltoall_perf
# nThread 1 nGpus 8 minBytes 536870912 maxBytes 1073741824 step: 52428800(bytes) warmup iters: 1 iters: 20 agg iters: 1
validation: 1 graph: 0
#
# Using devices
# Rank 0 Group 0 Pid 472216 on ps device 0 [0000:43:00] AMD Radeon AI PRO R9700
# Rank 1 Group 0 Pid 472216 on ps device 1 [0000:26:00] AMD Radeon AI PRO R9700
# Rank 2 Group 0 Pid 472216 on ps device 2 [0000:03:00] AMD Radeon AI PRO R9700
# Rank 3 Group 0 Pid 472216 on ps device 3 [0000:e3:00] AMD Radeon AI PRO R9700
# Rank 4 Group 0 Pid 472216 on ps device 4 [0000:c3:00] AMD Radeon AI PRO R9700
# Rank 5 Group 0 Pid 472216 on ps device 5 [0000:c6:00] AMD Radeon AI PRO R9700
# Rank 6 Group 0 Pid 472216 on ps device 6 [0000:a3:00] AMD Radeon AI PRO R9700
# Rank 7 Group 0 Pid 472216 on ps device 7 [0000:83:00] AMD Radeon AI PRO R9700
#
#
# out-of-place in-place
# size count type redop root time algbw busbw #wrong time algbw busbw #wrong
# (B) (elements) (us) (GB/s) (GB/s) (us) (GB/s) (GB/s)
# 536870912 16777216 float none -1 88073.9 6.10 5.33 0 36240.6 14.81 12.96 N/A
# 589299712 18415616 float none -1 54492.9 10.81 9.46 0 39774.2 14.82 12.96 N/A
# 641728512 20054016 float none -1 72488.0 8.85 7.75 0 43242.7 14.84 12.99 N/A
# 694157312 21692416 float none -1 72789.4 9.54 8.34 0 47536.9 14.60 12.78 N/A
# 746586112 23330816 float none -1 109303 6.83 5.98 0 50244.9 14.86 13.00 N/A
# 799014912 24969216 float none -1 56521.2 14.14 12.37 0 53737.5 14.87 13.01 N/A
# 851443712 26607616 float none -1 75840.4 11.23 9.82 0 58137.3 14.65 12.81 N/A
# 903872512 28246016 float none -1 81277.1 11.12 9.73 0 61558.8 14.68 12.85 N/A
# 956301312 29884416 float none -1 95495.0 10.01 8.76 0 65132.5 14.68 12.85 N/A
# 1008730112 31522816 float none -1 147495 6.84 5.98 0 67704.1 14.90 13.04 N/A
# 1061158912 33161216 float none -1 90040.4 11.79 10.31 0 71215.6 14.90 13.04 N/A
# Out of bounds values : 0 OK
# Avg bus bandwidth : 10.7332
#
# Collective test concluded: alltoall_perf
#
```

使用命令：./scatter_perf -g 8 -b 512M -e 1G -i 50m

```
# rccl-tests version 2.17.9-: rccl-headers=22707 rccl-library=22707
# Collective test starting: scatter_perf
# nThread 1 nGpus 8 minBytes 536870912 maxBytes 1073741824 step: 52428800(bytes) warmup iters: 1 iters: 20 agg iters: 1
validation: 1 graph: 0
#
# Using devices
# Rank 0 Group 0 Pid 472464 on ps device 0 [0000:43:00] AMD Radeon AI PRO R9700
# Rank 1 Group 0 Pid 472464 on ps device 1 [0000:26:00] AMD Radeon AI PRO R9700
# Rank 2 Group 0 Pid 472464 on ps device 2 [0000:03:00] AMD Radeon AI PRO R9700
# Rank 3 Group 0 Pid 472464 on ps device 3 [0000:e3:00] AMD Radeon AI PRO R9700
# Rank 4 Group 0 Pid 472464 on ps device 4 [0000:c3:00] AMD Radeon AI PRO R9700
# Rank 5 Group 0 Pid 472464 on ps device 5 [0000:c6:00] AMD Radeon AI PRO R9700
# Rank 6 Group 0 Pid 472464 on ps device 6 [0000:a3:00] AMD Radeon AI PRO R9700
# Rank 7 Group 0 Pid 472464 on ps device 7 [0000:83:00] AMD Radeon AI PRO R9700
#
#
# out-of-place in-place
# size count type redop root time algbw busbw #wrong time algbw busbw #wrong
# (B) (elements) (us) (GB/s) (GB/s) (us) (GB/s) (GB/s)
# 536870912 16777216 float none 0 24061.1 22.31 19.52 0 23860.0 22.50 19.69 0
# 589299712 18415616 float none 0 26604.5 22.15 19.38 0 26232.4 22.46 19.66 0
# 641728512 20054016 float none 0 29871.8 21.48 18.80 0 28577.9 22.46 19.65 0
# 694157312 21692416 float none 0 31385.8 22.12 19.35 0 30938.2 22.44 19.63 0
# 746586112 23330816 float none 0 33659.3 22.18 19.41 0 33457.4 22.31 19.53 0
# 799014912 24969216 float none 0 36459.1 21.92 19.18 0 35855.8 22.28 19.50 0
# 851443712 26607616 float none 0 38699.2 22.00 19.25 0 38161.7 22.31 19.52 0
# 903872512 28246016 float none 0 41019.6 22.04 19.28 0 40587.2 22.27 19.49 0
# 956301312 29884416 float none 0 43772.5 21.85 19.12 0 43091.6 22.19 19.42 0
# 1008730112 31522816 float none 0 45908.2 21.97 19.23 0 45324.4 22.26 19.47 0
# 1061158912 33161216 float none 0 48172.2 22.03 19.27 0 47629.6 22.28 19.49 0
# Out of bounds values : 0 OK
# Avg bus bandwidth : 19.4015
#
# Collective test concluded: scatter_perf
#
```

七、拓扑展示

使用命令：`rocm-smi --showtopo`

```

===== ROCm System Management Interface =====
===== Weight between two GPUs =====
      GPU0      GPU1      GPU2      GPU3      GPU4      GPU5      GPU6      GPU7
GPU0  0          52        52        72        72        72        72        72
GPU1  52        0          52        72        72        72        72        72
GPU2  52        52        0          72        72        72        72        72
GPU3  72        72        72        0          52        52        52        52
GPU4  72        72        72        52        0          52        52        52
GPU5  72        72        72        52        52        0          40        52
GPU6  72        72        72        52        52        40        0          52
GPU7  72        72        72        52        52        52        52        0

===== Hops between two GPUs =====
      GPU0      GPU1      GPU2      GPU3      GPU4      GPU5      GPU6      GPU7
GPU0  0          3          3          3          3          3          3          3
GPU1  3          0          3          3          3          3          3          3
GPU2  3          3          0          3          3          3          3          3
GPU3  3          3          3          0          3          3          3          3
GPU4  3          3          3          3          0          3          3          3
GPU5  3          3          3          3          3          0          2          3
GPU6  3          3          3          3          3          2          0          3
GPU7  3          3          3          3          3          3          3          0

===== Link Type between two GPUs =====
      GPU0      GPU1      GPU2      GPU3      GPU4      GPU5      GPU6      GPU7
GPU0  0          PCIE      PCIE      PCIE      PCIE      PCIE      PCIE      PCIE
GPU1  PCIE      0          PCIE      PCIE      PCIE      PCIE      PCIE      PCIE
GPU2  PCIE      PCIE      0          PCIE      PCIE      PCIE      PCIE      PCIE
GPU3  PCIE      PCIE      PCIE      0          PCIE      PCIE      PCIE      PCIE
GPU4  PCIE      PCIE      PCIE      PCIE      0          PCIE      PCIE      PCIE
GPU5  PCIE      PCIE      PCIE      PCIE      PCIE      0          PCIE      PCIE
GPU6  PCIE      PCIE      PCIE      PCIE      PCIE      PCIE      0          PCIE
GPU7  PCIE      PCIE      PCIE      PCIE      PCIE      PCIE      PCIE      0

===== Numa Nodes =====
GPU[0]      : (Topology) Numa Node: 3
GPU[0]      : (Topology) Numa Affinity: 3
GPU[1]      : (Topology) Numa Node: 2
GPU[1]      : (Topology) Numa Affinity: 2
GPU[2]      : (Topology) Numa Node: 1
GPU[2]      : (Topology) Numa Affinity: 1
GPU[3]      : (Topology) Numa Node: 7
GPU[3]      : (Topology) Numa Affinity: 7
GPU[4]      : (Topology) Numa Node: 6
GPU[4]      : (Topology) Numa Affinity: 6
GPU[5]      : (Topology) Numa Node: 5
GPU[5]      : (Topology) Numa Affinity: 5
GPU[6]      : (Topology) Numa Node: 5
GPU[6]      : (Topology) Numa Affinity: 5
GPU[7]      : (Topology) Numa Node: 4
GPU[7]      : (Topology) Numa Affinity: 4
===== End of ROCm SMI Log =====

```

八、Slurm 调用显卡测试

Slurm 版本：26.11

使用命令启动脚本：srun --gres=gpu:4 python main.py -a resnet18 --dummy

结论：使用 slurm 集群申请指定数量显卡，并运行大模型多卡训练，测试结果正常，无报错。amd r9700 可以搭配 slurm 集群使用。

```
srun --gres=gpu:4 python main.py -a resnet18 --dummy
/home/ps/.venv/lib/python3.12/site-packages/torch/utils/data/dataloader.py:424: UserWarning: This DataLoader will create 4
worker processes in total. Our suggested max number of worker in current system is 2, which is smaller than what this DataLoader
is going to create. Please be aware that excessive worker creation might get DataLoader running slow or even freeze, lower the
worker number to avoid potential slowness/freeze if necessary.
  self.check_worker_number_rationality()
/home/ps/.venv/lib/python3.12/site-packages/torch/utils/data/dataloader.py:432: UserWarning: This DataLoader will create 4
worker processes in total. Our suggested max number of worker in current system is 2, which is smaller than what this DataLoader
is going to create. Please be aware that excessive worker creation might get DataLoader running slow or even freeze, lower the
worker number to avoid potential slowness/freeze if necessary.
  self.check_worker_number_rationality()
Using device: cuda
=> creating model 'resnet18'
=> Dummy data is used!

Epoch: [0][ 1/5005] Time 50.514 (50.514) Data 3.571 ( 3.571) Loss 6.9916e+00 (6.9916e+00) Acc@1 0.39 ( 0.39)
Acc@5 0.39 ( 0.39)
Epoch: [0][ 11/5005] Time 1.412 ( 4.916) Data 1.129 ( 0.428) Loss 7.1556e+00 (7.0644e+00) Acc@1 0.00 ( 0.07)
Acc@5 0.78 ( 0.39)
Epoch: [0][ 21/5005] Time 0.574 ( 2.810) Data 0.373 ( 0.353) Loss 7.2407e+00 (7.1167e+00) Acc@1 0.00 ( 0.06)
Acc@5 0.00 ( 0.41)
Epoch: [0][ 31/5005] Time 0.472 ( 2.091) Data 0.297 ( 0.363) Loss 7.1569e+00 (7.1477e+00) Acc@1 0.39 ( 0.09)
Acc@5 0.78 ( 0.47)
Epoch: [0][ 41/5005] Time 1.808 ( 1.735) Data 1.574 ( 0.381) Loss 7.1126e+00 (7.1620e+00) Acc@1 0.39 ( 0.10)
Acc@5 0.78 ( 0.48)
Epoch: [0][ 51/5005] Time 0.240 ( 1.494) Data 0.000 ( 0.364) Loss 7.0660e+00 (7.1566e+00) Acc@1 0.00 ( 0.10)
Acc@5 0.39 ( 0.47)
Epoch: [0][ 61/5005] Time 1.370 ( 1.350) Data 1.176 ( 0.364) Loss 6.9877e+00 (7.1373e+00) Acc@1 0.00 ( 0.10)
Acc@5 1.17 ( 0.49)
Epoch: [0][ 71/5005] Time 0.157 ( 1.229) Data 0.003 ( 0.350) Loss 7.0056e+00 (7.1183e+00) Acc@1 0.00 ( 0.10)
Acc@5 0.00 ( 0.48)
Epoch: [0][ 81/5005] Time 1.536 ( 1.153) Data 1.357 ( 0.357) Loss 6.9959e+00 (7.0994e+00) Acc@1 0.00 ( 0.09)
Acc@5 0.39 ( 0.49)
Epoch: [0][ 91/5005] Time 0.227 ( 1.079) Data 0.000 ( 0.347) Loss 6.9545e+00 (7.0837e+00) Acc@1 1.56 ( 0.11)
Acc@5 2.34 ( 0.52)
Epoch: [0][ 101/5005] Time 1.536 ( 1.033) Data 1.268 ( 0.351) Loss 6.9667e+00 (7.0709e+00) Acc@1 0.00 ( 0.12)
Acc@5 0.00 ( 0.53)
Epoch: [0][ 111/5005] Time 0.143 ( 0.982) Data 0.000 ( 0.345) Loss 6.9480e+00 (7.0600e+00) Acc@1 0.00 ( 0.12)
Acc@5 1.17 ( 0.53)
Epoch: [0][ 121/5005] Time 1.533 ( 0.952) Data 1.339 ( 0.348) Loss 6.9186e+00 (7.0503e+00) Acc@1 0.00 ( 0.12)
Acc@5 0.39 ( 0.55)
Epoch: [0][ 131/5005] Time 0.159 ( 0.915) Data 0.000 ( 0.342) Loss 6.9494e+00 (7.0425e+00) Acc@1 0.00 ( 0.12)
Acc@5 0.39 ( 0.52)
Epoch: [0][ 141/5005] Time 1.496 ( 0.893) Data 1.231 ( 0.345) Loss 6.9497e+00 (7.0351e+00) Acc@1 0.00 ( 0.11)
Acc@5 0.39 ( 0.51)
Epoch: [0][ 151/5005] Time 0.178 ( 0.864) Data 0.004 ( 0.340) Loss 6.9190e+00 (7.0281e+00) Acc@1 0.39 ( 0.12)
Acc@5 0.78 ( 0.52)
Epoch: [0][ 161/5005] Time 1.547 ( 0.850) Data 1.348 ( 0.347) Loss 6.9274e+00 (7.0225e+00) Acc@1 0.00 ( 0.11)
Acc@5 0.78 ( 0.51)
Epoch: [0][ 171/5005] Time 0.174 ( 0.828) Data 0.000 ( 0.343) Loss 6.9466e+00 (7.0173e+00) Acc@1 0.39 ( 0.13)
Acc@5 0.39 ( 0.52)
```


ps: Thu May 28 17:05:38 2026

===== ROCm System Management Interface =====															
===== Concise Info =====															
Device	Node	IDS		Temp	Power	Partitions			SCLK	MCLK	Fan	Perf	PwrCap	VRAM%	GPU%
		(DID,		(GUID)	(Edge)	(Avg)	(Mem,	Compute,	ID)						
=====															
0	3	0x7551,	12518	42.0°C	97.0W	N/A,	N/A,	0	3254Mhz	1258Mhz	49.8%	auto	300.0W	9%	100%
1	2	0x7551,	47932	44.0°C	101.0W	N/A,	N/A,	0	3371Mhz	1258Mhz	49.8%	auto	300.0W	9%	100%
2	1	0x7551,	49367	40.0°C	96.0W	N/A,	N/A,	0	3156Mhz	1258Mhz	49.8%	auto	300.0W	9%	100%
3	8	0x7551,	36999	26.0°C	15.0W	N/A,	N/A,	0	0Mhz	96Mhz	49.8%	auto	300.0W	0%	0%
4	7	0x7551,	18590	26.0°C	13.0W	N/A,	N/A,	0	0Mhz	96Mhz	49.8%	auto	300.0W	0%	0%
5	5	0x7551,	24758	27.0°C	12.0W	N/A,	N/A,	0	0Mhz	96Mhz	49.8%	auto	300.0W	0%	0%
6	6	0x7551,	13173	24.0°C	12.0W	N/A,	N/A,	0	0Mhz	96Mhz	49.8%	auto	300.0W	0%	0%
7	4	0x7551,	47279	47.0°C	106.0W	N/A,	N/A,	0	3351Mhz	1258Mhz	49.8%	auto	300.0W	9%	100%
=====															
===== End of ROCm SMI Log =====															